



HarmBench Safety Evaluation — 49 Configurations, 8 Experiments, Llama Guard 3-1B Judge


(a) Exp 1: HarmBench — Qwen 2.5 / Llama / Gemma 2




(b) Exp 1: HarmBench — Gemma 3 / Qwen 3 / Others




(c) Exp 2: Quantization (Qwen 2.5-3B)




(d) Exp 3: Base vs. Instruct (Unaligned)




(e) Exp 4: Multilingual Safety Score (%)



(f) Exp 5: Jailbreak Resistance



(g) Exp 6: Multi-Turn Crescendo Attack



(h) Exp 7: Over-Refusal (JBB Benign)

